

class 07

Ruofan Kang (A17236920)


2023-10-24

Clustering

we will start today's lab with clustering methods, in particular so-called K-means. The main function for this in R is `kmeans()`

Let's try it on some made up data where we know that what the answer should be.

```
x <- rnorm(10000, mean=3)
hist(x)
```



60 points

```
tmp <- c(rnorm(30, mean=3), rnorm(30, mean=-3))
x <- cbind(x=tmp, y=rev(tmp))
x
```

```
##           x           y
```

```


## [1,] 3.240976 -2.766022
## [2,] 3.991392 -3.969736
## [3,] 2.163052 -3.104976
## [4,] 1.898661 -2.189034
## [5,] 3.976497 -4.582273
## [6,] 3.130927 -2.891582
## [7,] 5.168612 -3.659319
## [8,] 3.391826 -3.374629
## [9,] 3.210018 -2.129573
## [10,] 3.707352 -2.741579
## [11,] 3.409804 -2.395693
## [12,] 3.642331 -2.757303
## [13,] 2.735407 -3.775581
## [14,] 3.767412 -4.764453
## [15,] 4.460710 -5.233132
## [16,] 3.037118 -2.655066
## [17,] 3.327671 -2.233722
## [18,] 2.721973 -4.191819
## [19,] 3.439150 -1.636162
## [20,] 2.964130 -2.773621
## [21,] 1.942545 -2.565032
## [22,] 2.949764 -3.030119
## [23,] 2.747292 -3.179566
## [24,] 5.892171 -3.509852
## [25,] 2.890235 -1.899102
## [26,] 1.643839 -5.903929
## [27,] 2.071449 -2.813516
## [28,] 2.500929 -2.916589
## [29,] 4.002452 -2.268015
## [30,] 2.221102 -2.070330
## [31,] -2.070330 2.221102
## [32,] -2.268015 4.002452
## [33,] -2.916589 2.500929
## [34,] -2.813516 2.071449
## [35,] -5.903929 1.643839
## [36,] -1.899102 2.890235
## [37,] -3.509852 5.892171
## [38,] -3.179566 2.747292
## [39,] -3.030119 2.949764
## [40,] -2.565032 1.942545
## [41,] -2.773621 2.964130
## [42,] -1.636162 3.439150
## [43,] -4.191819 2.721973
## [44,] -2.233722 3.327671
## [45,] -2.655066 3.037118
## [46,] -5.233132 4.460710
## [47,] -4.764453 3.767412
## [48,] -3.775581 2.735407
## [49,] -2.757303 3.642331
## [50,] -2.395693 3.409804
## [51,] -2.741579 3.707352
## [52,] -2.129573 3.210018
## [53,] -3.374629 3.391826
## [54,] -3.659319 5.168612

```

```
## [55,] -2.891582  3.130927
## [56,] -4.582273  3.976497
## [57,] -2.189034  1.898661
## [58,] -3.104976  2.163052
## [59,] -3.969736  3.991392
## [60,] -2.766022  3.240976
```

We can pass this to the base R `plot()` function for a quick

```
plot(x)
```



```
k <- kmeans(x, centers=2, nstart=20)
k
```

```
## K-means clustering with 2 clusters of sizes 30, 30
##
## Cluster means:
##      x      y
## 1 -3.132711  3.208227
## 2  3.208227 -3.132711
##
## Clustering vector:
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
## [39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
##
## Within cluster sum of squares by cluster:
## [1] 55.68467 55.68467
## (between_SS / total_SS =  91.5 %)
```

```
##
## Available components:
##
## [1] "cluster"      "centers"      "totss"      "withinss"    "tot.withinss"
## [6] "betweenss"    "size"        "iter"      "ifault"      "
```

Q1. How many points are in each cluster

```
k$size
```

```
## [1] 30 30
```

Q. How we do get to the cluster membership/assignment.

```
k$cluster
```

```
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
## [39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q3. What about cluster centers?

```
k$centers
```

```
##          x          y
## 1 -3.132711  3.208227
## 2  3.208227 -3.132711
```


Now we got to the main results let's use them to plot our data with the kmeans result.

```
plot(x, col=k$cluster)
```




Q4. Plot x colored by the kmeans cluster assignment and add cluster centers as blue points

```
plot(x, col=k$cluster, pch=16)
```



Q5. Cluster the data again with `kmeans()` into 4 groups and plot the results.

```
k4 <- kmeans(x, center= 4, nstart=20)  
plot(x, col=k4$cluster, pch=16)
```



K-means is very popular mostly because it is fast and relatively straight forward to run and understand. It has a big limitation in that you need to tell it how many groups (k, or centers) you want.

#Hierarchical clustering

The main function in base R is called ‘`hclust()`’. You have to pass it in a “distance matrix” not just your input data.


you can generate a distance matrix with the “`dist()`”

```
hc <- hclust( dist(x))
hc
```

```
##
## Call:
## hclust(d = dist(x))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

```
plot(hc)
```


Cluster Dendrogram



To find the cluster(cluster membership vector)from a 'hclust()' result we can “cut” the tree at a certain height

```
plot(hc)
abline(h=8,col="red")
```

Cluster Dendrogram



```
grps <- cutree(hc, h=8)
```

```
table(grps)
```

```
## grps
## 1 2
## 30 30
```

Q6. Plot our hclust results.

Principal Component Analysis

PCA of UK food data

Read data showing the consumption in grams (per person, per week) of 17 different types of food-stuff measured and averaged in the four countries of the United kingdom.

Let's see how PCA can help us but first we can try conventional analysis.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
x
```

```
##           X England Wales Scotland N.Ireland
## 1           Cheese      105    103      103      66
## 2      Carcass_meat     245    227      242     267
## 3         Other_meat     685    803      750     586
## 4             Fish     147    160      122      93
```



```
## 5      Fats_and_oils      193  235      184      209
## 6          Sugars      156  175      147      139
## 7    Fresh_potatoes      720  874      566     1033
## 8          Fresh_Veg      253  265      171      143
## 9          Other_Veg      488  570      418      355
## 10 Processed_potatoes      198  203      220      187
## 11    Processed_Veg      360  365      337      334
## 12    Fresh_fruit     1102 1137      957      674
## 13          Cereals     1472 1582     1462     1494
## 14          Beverages       57   73       53       47
## 15    Soft_drinks     1374 1256     1572     1506
## 16  Alcoholic_drinks      375  475      458      135
## 17    Confectionery       54   64       62       41
```

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
rownames(x) <- x[,1]
x <- x[, -1]
head(x)
```

```
##              England Wales Scotland N.Ireland
## Cheese          105   103      103         66
## Carcass_meat     245  227      242        267
## Other_meat       685  803      750        586
## Fish             147  160      122         93
## Fats_and_oils     193  235      184        209
## Sugars           156  175      147        139
```

Q1.How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
## Complete the following code to find out how many rows and columns are in x?
dim(x)
```

```
## [1] 17  4
```

```
rownames(x) <- x[,1]
x <- x[, -1]
head(x)
```

```
##      Wales Scotland N.Ireland
## 105   103      103         66
## 245   227      242        267
## 685   803      750        586
## 147   160      122         93
## 193   235      184        209
## 156   175      147        139
```

```
dim(x)
```

```
## [1] 17  3
```


```
x <- read.csv(url, row.names=1)
head(x)
```

```
##              England Wales Scotland N.Ireland
## Cheese          105   103      103         66
## Carcass_meat     245  227      242        267
## Other_meat       685  803      750        586
```

```
## Fish      147  160  122   93
## Fats_and_oils 193  235  184  209
## Sugars    156  175  147  139
```


Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




Q3.Changing what optional argument in the above barplot() function results in the following plot?

```
barplot(as.matrix(x), beside=FALSE, col=rainbow(nrow(x)))
```



Q5. Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



Principal Component Analysis(PCA)

PCA can help us make sense of these types of datasets. Let's see how it works.

The main function in "base" R is called 'prcomp()'. In this case we want to first take the transpose of our input 'x' so the columns are the food types and the countries are the rows.

```
head( t(x) )
```

```
##      Cheese Carcass_meat Other_meat Fish Fats_and_oils Sugars
## England      105         245      685  147             193   156
## Wales        103         227      803  160             235   175
## Scotland     103         242      750  122             184   147
## N.Ireland      66         267      586   93             209   139
##
##      Fresh_potatoes Fresh_Veg Other_Veg Processed_potatoes
## England           720      253      488             198
## Wales             874      265      570             203
## Scotland          566      171      418             220
## N.Ireland         1033      143      355             187
##
##      Processed_Veg Fresh_fruit Cereals Beverages Soft_drinks
## England           360      1102     1472         57     1374
## Wales             365      1137     1582         73     1256
## Scotland          337      957     1462         53     1572
## N.Ireland          334      674     1494         47     1506
##
##      Alcoholic_drinks Confectionery
## England              375          54
## Wales                475          64
```

```
## Scotland          458          62
## N.Ireland         135          41
```

```
pca <- prcomp( t(x) )
summary(pca)
```


```
## Importance of components:
```

```
##              PC1      PC2      PC3      PC4
## Standard deviation 324.1502 212.7478 73.87622 2.921e-14
## Proportion of Variance 0.6744 0.2905 0.03503 0.000e+00
## Cumulative Proportion 0.6744 0.9650 1.00000 1.000e+00
```

```
pca$x
```


```
##              PC1      PC2      PC3      PC4
## England    -144.99315  -2.532999 105.768945 -9.152022e-15
## Wales      -240.52915 -224.646925 -56.475555  5.560040e-13
## Scotland   -91.86934  286.081786 -44.415495 -6.638419e-13
## N.Ireland  477.39164  -58.901862  -4.877895  1.329771e-13
```

```
plot( pca$x[,1], pca$x[,2], col=c("orange", "red", "blue", "darkgreen"),pch=16)
```



The “loadings” tells us how much the original variables (in our case the foods) contribute to the new variables i.e. the PCs.

```
## Lets focus on PC1 as it accounts for > 90% of variance
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```



```
pca$rotation
```

```
##           PC1           PC2           PC3           PC4
## Cheese      -0.056955380  0.016012850  0.02394295 -0.409382587
## Carcass_meat  0.047927628  0.013915823  0.06367111  0.729481922
## Other_meat   -0.258916658 -0.015331138 -0.55384854  0.331001134
## Fish        -0.084414983 -0.050754947  0.03906481  0.022375878
## Fats_and_oils -0.005193623 -0.095388656 -0.12522257  0.034512161
## Sugars       -0.037620983 -0.043021699 -0.03605745  0.024943337
## Fresh_potatoes 0.401402060 -0.715017078 -0.20668248  0.021396007
## Fresh_Veg    -0.151849942 -0.144900268  0.21382237  0.001606882
## Other_Veg    -0.243593729 -0.225450923 -0.05332841  0.031153231
## Processed_potatoes -0.026886233  0.042850761 -0.07364902 -0.017379680
## Processed_Veg -0.036488269 -0.045451802  0.05289191  0.021250980
## Fresh_fruit  -0.632640898 -0.177740743  0.40012865  0.227657348
## Cereals      -0.047702858 -0.212599678 -0.35884921  0.100043319
## Beverages    -0.026187756 -0.030560542 -0.04135860 -0.018382072
## Soft_drinks   0.232244140  0.555124311 -0.16942648  0.222319484
## Alcoholic_drinks -0.463968168  0.113536523 -0.49858320 -0.273126013
## Confectionery -0.029650201  0.005949921 -0.05232164  0.001890737
```

```
## The inbuilt biplot() can be useful for small datasets
```

```
biplot(pca)
```

